

Churn Model

April 30, 2022

0.1 Churn Prediction

Build a basic churn prediction model based on this data – predict the probability of customer churn in the next 6 month from the latest date. The approach would matter here, not the accuracy. You should clearly explain what are the data sanity checks you'll do, what parameters will you use to compare the models, how to measure the accuracy of the model – basically the entire end to end model cycle.

0.1.1

It is difficult to acquire new customers than to retain existing ones.

```
[248]: import pandas as pd
import numpy as np

import os
import datetime as dt
import utility_functions as uf

#Plotting related libraries
import matplotlib.pyplot as plt

#Modelling related libraries
from sklearn.linear_model import LogisticRegression as LR
from sklearn.model_selection import train_test_split
import joblib

#Model Evaluation Libraries
from sklearn.metrics import roc_curve, roc_auc_score
from sklearn.metrics import confusion_matrix, f1_score

[4]: #Read the UK related clean data
uk_data_NA_free_cust = pd.read_csv(os.getcwd() + '/Statistics/UK_data_NA_free.
→csv')

#Unique customers = 3,950
uk_data_NA_free_cust.shape #(3,61,878 X 14)
```

```
[4]: (361878, 14)
```

```
[5]: RFM_data = pd.read_csv(os.getcwd() + '/Final_Data.csv')
if('Unnamed: 0' in RFM_data.columns):
    RFM_data = RFM_data.drop('Unnamed: 0', axis=1)
RFM_data.head()
```

```
[5]:
```

	CustomerID	Days_SinceLastTxn	InvoiceNo	Counts	Total	Amount	R_Rank	\
0	13113.0	3794		40	10510.00	16.0		
1	15804.0	3794		19	3848.55	16.0		
2	13777.0	3794		41	25748.35	16.0		
3	17581.0	3794		31	10736.11	16.0		
4	12748.0	3794		224	29072.10	16.0		

	F_Rank	M_Rank	R_Rank_nrm	R_Rank_score	F_Rank_nrm	F_Rank_score	\
0	35.0	71.0	0.0	0	0.010133	0	
1	149.0	311.0	0.0	0	0.044107	0	
2	32.5	29.0	0.0	0	0.009388	0	
3	55.0	67.0	0.0	0	0.016093	0	
4	1.0	22.0	0.0	0	0.000000	0	

	M_Rank_nrm	M_Rank_score	RFM_score	Customer_Segments	Customer Tags
0	0.017726	0	0	2	Churned Customer
1	0.078501	0	0	2	Churned Customer
2	0.007090	0	0	2	Churned Customer
3	0.016713	0	0	2	Churned Customer
4	0.005318	0	0	2	Churned Customer

```
[6]: #3950
RFM_data.shape
```

```
[6]: (3950, 16)
```

```
[7]: #Ignore the variables 'M_Rank_score', 'R_Rank_score', 'F_Rank_score'
RFM_data[['M_Rank_score', 'R_Rank_score', 'F_Rank_score']].describe()
```

```
[7]:
```

	M_Rank_score	R_Rank_score	F_Rank_score
count	3950.000000	3950.000000	3950.000000
mean	0.000253	0.003544	0.300759
std	0.015911	0.059436	0.458646
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000
75%	0.000000	0.000000	1.000000
max	1.000000	1.000000	1.000000

```
[8]: #Drop the variables ['M_Rank_score', 'R_Rank_score', 'F_Rank_score']
print(RFM_data.shape)
score_cols = ['M_Rank_score', 'R_Rank_score', 'F_Rank_score']
```

```

if(set(score_cols).intersection(set(RFM_data.columns)) == set(score_cols)):
    RFM_data = RFM_data.drop(score_cols, axis=1)

print(RFM_data.shape)

```

(3950, 16)

(3950, 13)

0.2 Target Tagging

```

[9]: customer_tags = {0: 'New Customer', 1: 'Best Customer', \
                      2: 'Churned Customer', 3: 'Risk of Leaving'}

RFM_data['Customer Tags'] = RFM_data['Customer_Segments'].map(customer_tags)

cust_seg_obj = uf.GetStats(RFM_data[['CustomerID', 'Customer Tags']])

cust_seg_obj.categorical_distribution('Customer Tags')

```

```

[9]:
      Category  Counts  Percentage
0    Best Customer   1153      29.19
3  Churned Customer   1123      28.43
1    New Customer    918      23.24
2  Risk of Leaving    756      19.14

```

```

[10]: rfm_stats_obj = uf.GetStats(RFM_data)
rfm_stats = rfm_stats_obj.driver()
rfm_stats

```

```

[10]:
      Features  Missing Value Count  Missing Value Percentage  \
0    CustomerID                    0                    0.0
1  Days_SinceLastTxn                0                    0.0
2    InvoiceNo Counts                0                    0.0
3    Total Amount                  0                    0.0
4        R_Rank                    0                    0.0
5        F_Rank                    0                    0.0
6        M_Rank                    0                    0.0
7    R_Rank_nrm                    0                    0.0
8    F_Rank_nrm                    0                    0.0
9    M_Rank_nrm                    0                    0.0
10   RFM_score                    0                    0.0
11  Customer_Segments              0                    0.0
12   Customer Tags                0                    0.0

      Data Types  count      mean      std      min      25%  \
0    float64    3950.0  15562.029367  1576.848325  12346.00  14208.250000

```

1	int64	3950.0	3885.323038	100.236848	3794.00	3810.000000
2	int64	3950.0	5.027089	8.717306	1.00	1.000000
3	float64	3950.0	1713.385669	6548.608224	-4287.63	282.255000
4	float64	3950.0	1975.500000	1140.333924	16.00	965.000000
5	float64	3950.0	1975.500000	1119.520204	1.00	1108.000000
6	float64	3950.0	1975.500000	1140.411104	1.00	988.250000
7	float64	3950.0	0.498918	0.290346	0.00	0.241630
8	float64	3950.0	0.588437	0.333637	0.00	0.329906
9	float64	3950.0	0.500000	0.288785	0.00	0.250000
10	int64	3950.0	0.304557	0.468458	0.00	0.000000
11	int64	3950.0	1.434684	1.045756	0.00	1.000000
12	object	NaN	NaN	NaN	NaN	NaN

	50%	75%	max	# Categories
0	15571.500000	16913.750000	18287.00	NaN
1	3844.000000	3937.000000	4167.00	NaN
2	3.000000	5.000000	224.00	NaN
3	627.060000	1521.782500	256438.49	NaN
4	1991.500000	2962.500000	3943.50	NaN
5	1801.500000	3356.500000	3356.50	NaN
6	1975.500000	2962.750000	3950.00	NaN
7	0.502992	0.750223	1.00	NaN
8	0.536582	1.000000	1.00	NaN
9	0.500000	0.750000	1.00	NaN
10	0.000000	1.000000	2.00	NaN
11	1.000000	2.000000	3.00	NaN
12	NaN	NaN	NaN	4.0

0.3 Independent Variables

1. Days_SinceLastTxn
2. InvoiceNo Counts
3. Total Amount

0.4 Dependent Variable

Customer Tags

```
[11]: independent_vars = ['Days_SinceLastTxn', 'InvoiceNo Counts', 'Total Amount']
      target_var = 'Customer Tags'
      RFM_data[['CustomerID'] + independent_vars + [target_var]].head()
```

```
[11]: CustomerID  Days_SinceLastTxn  InvoiceNo Counts  Total Amount \
0      13113.0           3794           40      10510.00
1      15804.0           3794           19       3848.55
2      13777.0           3794           41      25748.35
3      17581.0           3794           31      10736.11
```

4	12748.0	3794	224	29072.10
---	---------	------	-----	----------

```

Customer Tags
0 Churned Customer
1 Churned Customer
2 Churned Customer
3 Churned Customer
4 Churned Customer

```

0.5 Predict Churn Probability

Strategy:-

Present:- Churn tags for the customers (given the recent date's input)

Get their input details 6 Months before recent date. This will serve as input for the model. Customer Tags for the recent date will serve as output.

Training Data, Testing Data and Validation Data

```

[12]: #Read the recency file
recency = pd.read_csv(os.getcwd() + "/Statistics/Recency_Customers.csv")
recency.head()

```

```

[12]: Unnamed: 0  CustomerID  Days_SinceLastTxn  Years_SinceLastTxn  \
0           0      13113.0           3794           10.394521
1           1      15804.0           3794           10.394521
2           2      13777.0           3794           10.394521
3           3      17581.0           3794           10.394521
4           4      12748.0           3794           10.394521

```

```

Months_SinceLastTxn
0      124.734247
1      124.734247
2      124.734247
3      124.734247
4      124.734247

```

```

[13]: #Days_SinceLastTxn = Days(Today's Date - InvoiceDate_new)
def get_prior_date(Invoice_date, n_months=6):
    """
    Input:-
    Invoice_date:- Format is datetime
    n_months:- The number of months
    This function will get the date n_months before the Invoice_date date.
    """
    if(type(Invoice_date) != dt.datetime):
        Invoice_date = Invoice_date.to_pydatetime()

```

```

        calculated_days = n_months*30 + (1*n_months//2) - 1 #Alternate 30, 31 days
        ↪months
        prior_date = Invoice_date - dt.timedelta(days=calculated_days)

        #Ignore time present in the date
        prior_date = dt.datetime(prior_date.year, prior_date.month, prior_date.day)

        return prior_date

#Date before 6 months from today
get_prior_date(dt.datetime.now(), 6)

```

[13]: `datetime.datetime(2021, 10, 30, 0, 0)`

```

[14]: #Get the input fields 6 months before current date
uk_data_NA_free_cust.head(3)
#Months since last transaction
#Necessary Columns
'''
['CustomerID', 'Quantity', 'UnitPrice', ]
'''

```

[14]: `"\n['CustomerID', 'Quantity', 'UnitPrice',]\n"`

```

[15]: #6-Months prior date
customer_grp = uk_data_NA_free_cust.groupby('CustomerID') #CustomerID is the
        ↪index

```

0.6 R

```

[161]: class RFM_Prior_Date():
        '''
        Input:-
        df:- A set of records belonging to a particular customer
        Return:-
        R:- Compute fields related to Recency for the customer
        F:- Compute fields related to Frequency for the customer
        M:- Compute fields related to Monetary for the customer
        '''

        def __init__(self, df, months=6):
            self.data = df
            self.months_before = months

        def get_prior_date(self, Invoice_date, n_months=6):
            '''
            This function will get the date n_months before the Invoice_date date.
            Input:-
            '''

```

```

Invoice_date:- Format is datetime
n_months:- The number of months

Return:-
prior_date:- The date n_months before the Invoice_date date.
'''

if(type(Invoice_date) != dt.datetime):
    Invoice_date = Invoice_date.to_pydatetime()

calculated_days = n_months*30 + (1*n_months//2) - 1 #Alternate 30, 31
→days months
prior_date = Invoice_date - dt.timedelta(days=calculated_days)

#Ignore time present in the date
prior_date = dt.datetime(prior_date.year, prior_date.month, prior_date.
→day)

return prior_date

def get_prior_date_records(self, df_grp):

    '''
    Input:-
    df_grp:- The records corresponding to a particular customer

    Return:-
    prior_data:- The data corresponding to previous date
    prior_date_today:- The previous date before today's date
    The prior_date_today will be needed for date calculations in the past.
    '''

    global customers_with_prior_record #Add the customerID whose 6 months
→prior record exists

    df_grp['InvoiceDateExclusive'] = df_grp['InvoiceDateExclusive'].
→apply(pd.to_datetime)
    df_grp = df_grp.sort_values(by='InvoiceDate_new', ascending=False)

    #Create a month variable. Compare 6 months prior using this variable
    df_grp['Months'] = list(map(lambda invoice_dt:int(str(invoice_dt).
→split()[0].split('-')[1]),\
                                df_grp['InvoiceDateExclusive'].tolist()))

    df_grp = df_grp.reset_index()
    current_date = df_grp.loc[0, 'InvoiceDateExclusive']

    #Get the current date 6 months before today

```

[illegible]


```

        prior_recency = prior_recency.sort_values(by='TimeSinceLastTxn',
↪ascending=False)

        #Most recent transactions(InvoiceDate_new) will be at the top
        prior_recency = pd.DataFrame(prior_recency.iloc[0,:]).T #Latest
↪record

        prior_recency = prior_recency.reset_index()
        customers_with_prior_record.add(prior_recency.loc[0, 'CustomerID'])
↪#Add the customer to the set

        return prior_recency

#F
def get_frequency(self, df_grp):
    '''
    Input:-
    df_grp:- The records corresponding to a particular customer

    Return:-
    Record with 1 row
    prior_frequency:- The frequency count for that customer
    '''
    #Get previous date's data and count of Invoices then for each customer
    prior_frequency, prior_date_today = self.get_prior_date_records(df_grp)

    if(prior_frequency.shape[0]>0):
        #Count Unique Invoices for each customer
        unique_invoices_count = len(np.unique(prior_frequency['InvoiceNo']))

        prior_frequency = prior_frequency.reset_index()

        customer_id = prior_frequency.loc[0, 'CustomerID']

        prior_frequency = pd.DataFrame({'CustomerID':[customer_id], \
                                         'InvoiceNo Counts':
↪[unique_invoices_count]},\
                                         columns = ['CustomerID', 'InvoiceNo',
↪Counts'])

        return prior_frequency[['CustomerID', 'InvoiceNo Counts']]

#M
def get_monetary(self, df_grp):
    '''
    Input:-
    df_grp:- The records corresponding to a particular customer

```

```

Return:-
Record with 1 row
prior_monetary:- The total amount spent by the customer
'''

#Get previous date's data and count of Invoices then for each customer
prior_monetary, prior_date_today = self.get_prior_date_records(df_grp)
if(prior_monetary.shape[0]>0):
    total_amount_sum = sum(df_grp['Total Amount'])#, max(df['Total_
↪Amount']))))

    prior_monetary = prior_monetary.reset_index()
    customer_id = prior_monetary.loc[0, 'CustomerID']

    prior_monetary = pd.DataFrame({'CustomerID':[customer_id], \
                                  'Total Amount Sum':
↪[total_amount_sum]},\
                                  columns = ['CustomerID', 'Total_
↪Amount Sum'])

    return prior_monetary[['CustomerID', 'Total Amount Sum']]

def remove_unwanted_cols(self, data):
    #Remove unwanted columns created as a result of .reset_index()
    for unwanted_col in ['index', 'Unnamed: 0', 'level_0']:
        if(unwanted_col in data.columns):
            data = data.drop(unwanted_col, axis=1)

    if('CustomerID' in data.index):
        if('CustomerID' in data.columns):
            data = data.drop('CustomerID', axis=1)

    data = data.reset_index()

    #print(data)
    return data

def driver(self):
    #Group the data basis the CustomerID
    customer_grp = self.data.groupby('CustomerID')

    #Apply R,F and M on each of the groups
    prior_recency = customer_grp.apply(lambda df:self.get_recency(df))

    prior_frequency = customer_grp.apply(lambda df:self.get_frequency(df))

    prior_monetary = customer_grp.apply(lambda df:self.get_monetary(df))

```

```

#Remove unwanted columns created as a consequence of .reset_index
RFM_data_frames = [prior_recency, prior_frequency, prior_monetary]
prior_recency, prior_frequency, prior_monetary = list(map(
    self.
    ↪remove_unwanted_cols, RFM_data_frames))
    print("Data Shapes:- \nR:-{}\nF:-{}\nM:-{}" .format(prior_recency.
    ↪shape, \
                                                    prior_frequency.
    ↪shape, \
                                                    prior_monetary.
    ↪shape))
    return prior_recency, prior_frequency, prior_monetary
'''
#Merge the frequency and recency data
RFM = pd.merge(prior_recency, prior_frequency, on='CustomerID',
    ↪how='left')
    return RFM
'''

```

```

[162]: RFM_object = RFM_Prior_Date(uk_data_NA_free_cust)
prior_recency, prior_frequency, prior_monetary = RFM_object.driver()

```

Data Shapes:-

R:-(764, 18)

F:-(764, 2)

M:-(764, 2)

```

[172]: prior_monetary.head()

```

```

[172]:
CustomerID  Total Amount Sum
CustomerID
12747.0     0      12747.0      4196.01
12748.0     0      12748.0     29072.10
12823.0     0      12823.0      1759.50
12826.0     0      12826.0      1468.12
12839.0     0      12839.0      5583.62

```

```

[182]: from functools import reduce
'''

dfs = [df1, df2, df3, df4, df5, df6]
df_final = reduce(lambda left, right: pd.
    ↪merge(left, right, on='some_common_column_name'), dfs)

'''
def joining_RFM_data(R, F, M):
    def remove_unwanted_cols(data):

```

```

#data = data.reset_index()
#Remove unwanted columns created as a result of .reset_index()
if('CustomerID' in data.columns):
    data = data.drop('CustomerID', axis=1)
    data = data.reset_index()
    for unwanted_col in ['index', 'Unnamed: 0', 'level_0', 'level_1']:
        if(unwanted_col in data.columns):
            data = data.drop(unwanted_col, axis=1)

    return data

F = remove_unwanted_cols(prior_frequency)
R = remove_unwanted_cols(prior_recency)
M = remove_unwanted_cols(prior_monetary)

dfs = [R, F, M]
RFM = reduce(lambda left, right: pd.merge(left, right, on='CustomerID'),
→dfs)
return RFM

```

```
[187]: RFM_prior = joining_RFM_data(prior_recency, prior_frequency, prior_monetary)
```

```

#Write the Modelling data to a file
RFM_prior.to_csv(os.getcwd() + '/Modelling/modelling_data.csv')
RFM_prior.head()

```

```
[187]:
```

	CustomerID	InvoiceNo	StockCode	Description	Quantity	\
0	12747.0	558265	22424	ENAMEL BREAD BIN CREAM	1	
1	12748.0	555620	22436	12 COLOURED PARTY BALLOONS	20	
2	12823.0	548245	48138	DOORMAT UNION FLAG	60	
3	12826.0	556681	84077	WORLD WAR 2 GLIDERS ASSTD DESIGNS	48	
4	12839.0	556229	20712	JUMBO BAG WOODLAND ANIMALS	10	

	InvoiceDate	UnitPrice	Country	InvoiceNo_Date	\
0	6/28/2011 10:06	12.75	United Kingdom	558265_6/28/2011 10:06	
1	6/6/2011 11:46	0.65	United Kingdom	555620_6/6/2011 11:46	
2	3/30/2011 10:36	7.65	United Kingdom	548245_3/30/2011 10:36	
3	6/14/2011 9:41	0.29	United Kingdom	556681_6/14/2011 9:41	
4	6/9/2011 14:43	2.08	United Kingdom	556229_6/9/2011 14:43	

	InvoiceDateExclusive	InvoiceTimeExclusive	InvoiceDate_new	Total Amount	\
0	2011-06-28	10:06	2011-06-28 10:06:00	12.75	
1	2011-06-06	11:46	2011-06-06 11:46:00	13	
2	2011-03-30	10:36	2011-03-30 10:36:00	459	
3	2011-06-14	9:41	2011-06-14 09:41:00	13.92	
4	2011-06-09	14:43	2011-06-09 14:43:00	20.8	

	Months	TimeSinceLastTxn	Days_SinceLastTxn	Years_SinceLastTxn	\
0	6	3777 days	3777	10.3479	
1	6	3799 days	3799	10.4082	
2	3	3867 days	3867	10.5945	
3	6	3791 days	3791	10.3863	
4	6	3796 days	3796	10.4	

	Months_SinceLastTxn	InvoiceNo	Counts	Total	Amount	Sum
0	124.175		1	4196.01		
1	124.899		16	29072.10		
2	127.134		1	1759.50		
3	124.636		2	1468.12		
4	124.8		1	5583.62		

0.7 Modelling

Input is 6 months before the current date.

Target(Customer Tags) is current target present in RFM_data prepared during part 1.

```
[190]: #Map the target class which is 6 months ahead for each of the 764 customers
```

```
RFM_prior = pd.merge(RFM_prior, RFM_data[['CustomerID', target]], \
                      on='CustomerID', how='left')

#Target Distribution
modelling_data_stats_obj = uf.GetStats(RFM_prior)
modelling_data_stats_obj.categorical_distribution('Customer Tags')
```

```
[190]:
```

	Category	Counts	Percentage
3	Churned Customer	531	69.50
1	New Customer	187	24.48
2	Risk of Leaving	31	4.06
0	Best Customer	15	1.96

0.7.1 Modelling Data Statistics

```
[191]: modelling_data_stats_obj.driver()
```

```
[191]:
```

	Features	Missing Value Count	Missing Value Percentage	\
0	CustomerID	0	0.0	
1	InvoiceDate_new	0	0.0	
2	Total Amount Sum	0	0.0	
3	InvoiceNo Counts	0	0.0	
4	Months_SinceLastTxn	0	0.0	
5	Years_SinceLastTxn	0	0.0	

6	Days_SinceLastTxn	0	0.0
7	TimeSinceLastTxn	0	0.0
8	Months	0	0.0
9	Total Amount	0	0.0
10	InvoiceTimeExclusive	0	0.0
11	InvoiceNo	0	0.0
12	InvoiceDateExclusive	0	0.0
13	InvoiceNo_Date	0	0.0
14	Country	0	0.0
15	UnitPrice	0	0.0
16	InvoiceDate	0	0.0
17	Quantity	0	0.0
18	Description	0	0.0
19	StockCode	0	0.0
20	Customer Tags	0	0.0

	Data Types	count	mean	std	min	25%	\
0	float64	764.0	15541.700262	1581.888356	12747.00	14188.25	
1	object	NaN	NaN	NaN	NaN	NaN	
2	float64	764.0	4993.985825	13874.843717	-24.05	1054.14	
3	int64	764.0	1.723822	1.353955	1.00	1.00	
4	object	NaN	NaN	NaN	NaN	NaN	
5	object	NaN	NaN	NaN	NaN	NaN	
6	object	NaN	NaN	NaN	NaN	NaN	
7	timedelta64[ns]	NaN	NaN	NaN	NaN	NaN	
8	object	NaN	NaN	NaN	NaN	NaN	
9	object	NaN	NaN	NaN	NaN	NaN	
10	object	NaN	NaN	NaN	NaN	NaN	
11	object	NaN	NaN	NaN	NaN	NaN	
12	datetime64[ns]	NaN	NaN	NaN	NaN	NaN	
13	object	NaN	NaN	NaN	NaN	NaN	
14	object	NaN	NaN	NaN	NaN	NaN	
15	object	NaN	NaN	NaN	NaN	NaN	
16	object	NaN	NaN	NaN	NaN	NaN	
17	object	NaN	NaN	NaN	NaN	NaN	
18	object	NaN	NaN	NaN	NaN	NaN	
19	object	NaN	NaN	NaN	NaN	NaN	
20	object	NaN	NaN	NaN	NaN	NaN	

	50%	75%	max	# Categories
0	15539.500	16878.2500	18283.00	NaN
1	NaN	NaN	NaN	759.0
2	2202.315	4416.5525	256438.49	NaN
3	1.000	2.0000	16.00	NaN
4	NaN	NaN	NaN	131.0
5	NaN	NaN	NaN	131.0
6	NaN	NaN	NaN	131.0

7	NaN	NaN	NaN	NaN
8	NaN	NaN	NaN	7.0
9	NaN	NaN	NaN	285.0
10	NaN	NaN	NaN	404.0
11	NaN	NaN	NaN	764.0
12	NaN	NaN	NaN	NaN
13	NaN	NaN	NaN	764.0
14	NaN	NaN	NaN	1.0
15	NaN	NaN	NaN	80.0
16	NaN	NaN	NaN	759.0
17	NaN	NaN	NaN	56.0
18	NaN	NaN	NaN	513.0
19	NaN	NaN	NaN	509.0
20	NaN	NaN	NaN	4.0

```
[198]: #Convert the columns to their correct data types
for col in ['Months_SinceLastTxn', 'Years_SinceLastTxn', 'Days_SinceLastTxn']:
    RFM_prior[col] = pd.to_numeric(RFM_prior[col])

RFM_prior['InvoiceDate_new'] = pd.to_datetime(RFM_prior['InvoiceDate_new'])

RFM_prior[['Months_SinceLastTxn', 'Years_SinceLastTxn', \
            'Days_SinceLastTxn', 'InvoiceDate_new']].dtypes
```

```
[198]: Months_SinceLastTxn      float64
Years_SinceLastTxn          float64
Days_SinceLastTxn           int64
InvoiceDate_new             datetime64[ns]
dtype: object
```

```
[226]: independent_variables = ['Months_SinceLastTxn', 'Years_SinceLastTxn', \
                                'Days_SinceLastTxn', 'Total Amount Sum']
RFM_prior[['CustomerID'] + independent_variables + [target]].head()
```

```
[226]: CustomerID  Months_SinceLastTxn  Years_SinceLastTxn  Days_SinceLastTxn  \
0      12747.0      124.175342      10.347945      3777
1      12748.0      124.898630      10.408219      3799
2      12823.0      127.134247      10.594521      3867
3      12826.0      124.635616      10.386301      3791
4      12839.0      124.800000      10.400000      3796
```

	Total Amount Sum	Customer Tags
0	4196.01	Churned Customer
1	29072.10	Churned Customer
2	1759.50	New Customer
3	1468.12	Churned Customer
4	5583.62	Churned Customer

```
[227]: #Binary Classification --> Churn or not Churn
RFM_prior['Target'] = RFM_prior['Customer Tags'].apply(lambda target_class:1
↳if(target_class=='Churned Customer') \
else 0)
```

```
[228]: modelling_data_stats_obj = uf.GetStats(RFM_prior)
modelling_data_stats_obj.categorical_distribution('Target')
```

```
[228]:      Category  Counts  Percentage
1          1      531         69.5
0          0      233         30.5
```

```
[244]: #Split the data into train and test sets #764 Observations
#Train Size = , Test Size =
X_train, X_test, y_train, y_test = train_test_split(RFM_prior[['CustomerID'] +
↳independent_variables],\
                                                    RFM_prior['Target'],\
                                                    test_size = 0.15)

train_data = pd.concat([X_train, y_train], axis=1)
test_data = pd.concat([X_test, y_test], axis=1)

print("Train Data Shape:- {} \nTest Data Shape:- {}".format(train_data.shape, \
test_data.shape))
```

```
Train Data Shape:- (649, 6)
Test Data Shape:- (115, 6)
```

```
[230]: train_data.head()
```

```
[230]:      CustomerID  Months_SinceLastTxn  Years_SinceLastTxn  Days_SinceLastTxn  \
762      18260.0          130.553425          10.879452          3971
341      15199.0          126.509589          10.542466          3848
137      13798.0          124.898630          10.408219          3799
214      14336.0          125.983562          10.498630          3832
234      14507.0          125.852055          10.487671          3828

      Total Amount Sum  Customer Tags
762          2595.00    New Customer
341          1268.70    New Customer
137         36351.42  Churned Customer
214          1614.91  Churned Customer
234          1368.18  Churned Customer
```

```
[231]: test_data.head()
```



```
[231]:
```

	CustomerID	Months_SinceLastTxn	Years_SinceLastTxn	Days_SinceLastTxn	\
565	16837.0	126.016438	10.501370	3833	
331	15132.0	125.326027	10.443836	3812	
146	13871.0	125.556164	10.463014	3819	
362	15367.0	126.838356	10.569863	3858	
504	16515.0	126.871233	10.572603	3859	

	Total Amount Sum	Customer Tags
565	3167.73	Churned Customer
331	977.93	Churned Customer
146	6389.80	Churned Customer
362	1867.69	New Customer
504	1627.13	New Customer

```
[259]: class Modeling():

    def __init__(self, X_train, X_test, y_train, y_test, model,
name_of_classifier):
        self.X_train = X_train
        self.X_test = X_test
        self.y_train = y_train
        self.y_test = y_test
        self.model = model
        self.clf_name = name_of_classifier

    def auc_roc(self, train_test='train'):
        if(train_test!='train'):
            print("Testing Data")
            X_data = self.X_test
            y_data = self.y_test
        else:
            print("Training Data")
            X_data = self.X_train
            y_data = self.y_train

        predicted_probability = self.model.predict_proba(X_data)

        fpr, tpr, _ = roc_curve(y_data, predicted_probability[:,1])
        auc = roc_auc_score(y_data, predicted_probability[:,1])
        plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
        plt.legend(loc=4)
        plt.xlabel("False Positive Rate")
        plt.ylabel("True Positive Rate")
        plt.savefig("AUC " + self.clf_name + ".png")
        plt.show()

    def evaluate(self, CM, train_test='train'):
```

```

accuracy = CM.diagonal().sum()/CM.sum()
precision = CM[0][0]/(CM[0][0] + CM[0][1])
recall = CM[0][0]/(CM[0][0] + CM[1][0])
f_measure = (2*precision*recall)/(precision+recall)
if(train_test=='train'):
    print("Training Data")
else:
    print("Testing Data")

print(CM)
print("Accuracy:- {:.2f}\nPrecision:- {:.2f}\nRecall:- {:.2f}\nF-measure:- {:.2f}" \
      .format(accuracy*100, precision*100, recall*100, f_measure*100))

def prediction(self, f=1):
    #X_train, X_test, y_train, y_test = train_test_split(self.X, self.y,
    #test_size=test_size/100, random_state=42)

    #Fit the model
    model = self.model.fit(self.X_train, self.y_train)

    #Predict on test set
    predicted_train, predicted_test = model.predict(self.X_train), model.
    predict(self.X_test)

    #Results
    result_test = pd.DataFrame({'Actual':self.y_test.ravel(), 'Predicted':
    predicted_test.ravel()}, columns=['Actual', 'Predicted'])
    result_train = pd.DataFrame({'Actual':self.y_train.ravel(), 'Predicted':
    predicted_train.ravel()}, columns=['Actual', 'Predicted'])

    #Save the model
    joblib.dump(model, os.getcwd() + "/Modelling/" + self.clf_name + ".pkl")

    #Print the AUC_ROC curve
    self.auc_roc() #Train
    self.auc_roc('test') #Test

    #Evaluate the model
    CM_train = confusion_matrix(self.y_train.ravel(), predicted_train.
    ravel(), labels=None, sample_weight=None)
    self.evaluate(CM_train)

    CM_test = confusion_matrix(self.y_test.ravel(), predicted_test.ravel(),
    labels=None, sample_weight=None)
    self.evaluate(CM_test, "test")

```


```
[260]: lr = LR()

logistic_model = lr.fit(X_train[independent_variables], y_train.tolist())
test_prediction = logistic_model.predict(X_train[independent_variables])


train_prediction = logistic_model.predict(X_test[independent_variables])

[261]: logistic_model = Modeling(X_train, X_test, y_train, y_test, lr, 'Logistic_
↳Regression')
logistic_model.prediction()
```

Training Data



Testing Data



Training Data

[[146 53]

[59 391]]

Accuracy:- 82.74

Precision:- 73.37

Recall:- 71.22

F-measure:- 72.28

Testing Data

[[23 11]

[10 71]]

Accuracy:- 81.74

Precision:- 67.65

Recall:- 69.70

F-measure:- 68.66

[]:

[]:

[]: